

Scientific Computing

Project 2a

Neural Networks

Matthew Crean


February 8, 2026

Background

- ▶ A neuron is a simple scalar functional unit, given by
 $\sigma : \mathbb{R} \rightarrow \mathbb{R}$
- ▶ In a feed-forward neural network, neurons are arranged in layers
 - ▶ Each neuron in a layer is connected to every neuron in the next layer
 - ▶ Each connection has an associated weight **W**
 - ▶ Each neuron has an associated bias **b**
- ▶ A network is trained by specifying 'training data' given a set of inputs and desired outputs, which changes the weights and biases to minimise the cost function
 - ▶ The cost function quantifies the error of the network's output compared to the desired output from the training data
- ▶ New inputs can then be classified by the trained network to retrieve the expected output

Neural Network Example: (2,3,3,1)

- ▶ Activation vector at each layer is $\mathbf{a}^{[l]} = \sigma_l (\mathbf{W}^l \mathbf{a}^{[l-1]} + \mathbf{b}^{[l]})$
 - ▶ $l = 1, \dots, L$ corresponds to the layer
 - ▶ σ_l is the activation function for layer l
 - ▶ $\mathbf{W}^{[l]}$ and $\mathbf{b}^{[l]}$ are the weight matrix and bias vector of size $n_l \times n_{l-1}$ and length n_l respectively
 - ▶ $\mathbf{a}^{[l]} = (a_1^{[l]}, \dots, a_{n_l}^{[l]})^\top$, n_l is the number of neurons in layer l
 - ▶ $\mathbf{a}^{[0]} = (x_1, x_2)^\top$, the input to the network



Layer 0 (input)

Layer 1

Layer 2 Layer 3 (output)

Neural Network Implementation


- ▶ Task: implement a feed-forward neural network for binary classification
- ▶ NeuralNetwork class:
 - ▶ Represents the entire neural network and the flow of data through it
 - ▶ Stores:
 - ▶ number of input variables to network
 - ▶ vector of all non-input layers ('NeuralNetworkLayer' objects)
 - ▶ Uses '*feed_forward*' to compute network output
 - ▶ does not modify model parameters
- ▶ NeuralNetworkLayer class:
 - ▶ Represents a single non-input layer of the network
 - ▶ For each layer, stores:
 - ▶ weight matrix
 - ▶ bias vector
 - ▶ activation function
 - ▶ Uses '*forward*' function which takes input activation vector and computes next layer's activation vector (computes $a^{[l]}$)

Optimising Weights and Biases

- ▶ For the training algorithm, we aim to minimise the cost to find our weight matrices and bias vectors
- ▶ Derivatives of this function tell us how to change the network parameters to reduce the cost function
 - ▶ Derivatives applied practically using stochastic gradient descent (SGD), where parameters are updated using individual training examples chosen at random rather than the full dataset
- ▶ These derivatives must be calculated via either:
 - ▶ back-propagation
 - ▶ propagates gradients 'backward' through the network
 - ▶ finite differencing
 - ▶ approximates derivatives by perturbing each parameter individually and recomputing cost
- ▶ Implemented within '*train*' function, given training rate η
 - ▶ η acts as a step-size controlling how strongly gradients influence the update of weights and biases at each iteration
 - ▶ training stopped once either target cost τ reached or maximum number of iterations reached

Simple Test Case


- ▶ Build neural network with 4 layers containing (2,3,3,1) neurons respectively
 - ▶ Activation function: $\tanh(x)$
 - ▶ Learning rate: $\eta = 0.1$
 - ▶ Target cost: $\tau = 10^{-4}$
 - ▶ Max iterations: 10^5
 - ▶ Initialise weights and biases from $N(0, 0.1^2)$



Plot shows total cost decreases to below target cost within maximum number of iterations.


- ▶ Initial rapid decay
- ▶ Oscillations due to stochastic nature of SGD (point randomly selected from training set each iteration)
- ▶ Network trains successfully

Simple Test Case: Contour Plots



Contour plot of network output after training once



- ▶ Network has learned a decision boundary that separates the two classes accurately



Contour plot after repeating training four times





- ▶ Slight variations in boundary due to different random initialisations and stochastic training process
- ▶ Training process is very reliable

Spiral Test Case







- ▶ Build network on the spiral dataset, with a (2,4,4,1) structure
 - ▶ Iteration plot shows cost initially decreases, but then starts to oscillate around cost 0.35
 - ▶ Network fails to converge to target cost
 - ▶ Contour plot shows the network captures some of the spiral shape, but not very accurately
- ▶ The (2,4,4,1) network is not expressive enough to capture the spiral structure; a wider (increasing neurons) or deeper (increasing layers) network should be tested

Spiral: Deeper Network Plots



Spiral: Wider Network Plots



Spiral: Varying Network Structure Analysis


- ▶ Deeper networks:
 - ▶ Increasing the number of layers shows decrease in performance of the neural network
 - ▶ (2,4,4,4,1) structure provides worse approximation of boundary and more unpredictable oscillations
 - ▶ (2,4,4,4,4,1) structure fails completely; does not classify the points at all and total cost does not go below 0.5
 - ▶ For the spiral data, deeper networks are worse as noise from each update in SGD passes through more layers, so it has a more drastic impact on resulting network
- ▶ Wider networks:
 - ▶ Much more effective at classifying spiral data
 - ▶ Neither networks fully converge but still approximate boundary for classification better than initial structure
 - ▶ Contour plot shows that (2,16,16,1) structure approximates classification boundary very well, despite some outliers
 - ▶ Increasing neurons is more effective as network has more parameters per layer to be optimised, improving the approximation

Theoretical Computational Cost

- ▶ Aim: Assess cost of training algorithm
- ▶ Suppose a network consists of N_{layer} layers, each with N_{neuron} neurons (both large to neglect small order operations)
- ▶ Single feed-forward operation:
 - ▶ $\mathbf{a}^{[l]} = \sigma_l(\mathbf{W}^{[l]}\mathbf{a}^{[l-1]} + \mathbf{b}^{[l]})$
 - ▶ $\mathbf{W}^{[l]}$ $N_{\text{neuron}} \times N_{\text{neuron}}$ matrix and $\mathbf{a}^{[l-1]}$ length N_{neuron} vector
 - ▶ $\mathbf{W}^{[l]}\mathbf{a}^{[l-1]}$ requires $O(N_{\text{neuron}}^2)$ operations
 - ▶ Repeated over all layers, so feed-forward operation costs $O(N_{\text{layer}}N_{\text{neuron}}^2)$
- ▶ Derivatives by finite-differencing:
 - ▶ To compute derivatives by finite-differencing, we must complete a full feed-forward for each parameter
 - ▶ There are $O(N_{\text{layer}}N_{\text{neuron}}^2)$ parameters
 - ▶ One feed-forward \times number of parameters = $O(N_{\text{layer}}^2N_{\text{neuron}}^4)$
- ▶ Derivatives by back-propagation:
 - ▶ We pass forward once and pass backwards once
 - ▶ $2 \times O(N_{\text{layer}}N_{\text{neuron}}^2) = O(N_{\text{layer}}N_{\text{neuron}}^2)$

Computational Cost: Varying Neurons


- ▶ Investigate costs of training by fixing $N_{\text{layer}} = 10$ and varying $N_{\text{neuron}} \in \{10, 20, 40, 80, 160\}$
 - ▶ Simulate random input and output data, initialise weights and biases randomly, record time taken for a single operation
 - ▶ Log-log plot should show gradient of 2 for feed-forward and back-propagation, and 4 for finite differencing



- ▶ Produces approximate straight lines
- ▶ Taking start and end points, find gradient:
 - ▶ Feed-forward ≈ 1.44
 - ▶ Finite differencing ≈ 3.41
 - ▶ Back-propagation ≈ 1.51
- ▶ Measured gradients are lower than the asymptotic predictions (2 and 4) due to limited range of N_{neuron} to prevent long run-times
- ▶ Results align with theory

Computational Cost: Varying Layers

- ▶ Further investigate costs of training by fixing $N_{\text{neuron}} = 40$ and varying $N_{\text{layer}} \in \{2, 4, 8, 16, 32\}$
 - ▶ Same set-up as varying neurons
 - ▶ Log-log plot should show gradient of 1 for feed-forward and back-propagation, and 2 for finite differencing




- ▶ Similarly, taking start and endpoints, find gradient of:
 - ▶ Feed-forward ≈ 0.96
 - ▶ Finite differencing ≈ 2.40
 - ▶ Back-propagation ≈ 0.99
- ▶ Measured gradients for feed-forward and back-propagation align with theory
- ▶ Finite differencing may have larger than predicted gradient due to increased depth resulting in more storage required, so slower run time

Interlacing Computation

- ▶ Currently in the implemented training algorithm, we:
 - ▶ Store each partial differential of each component of weight matrix and bias vector
 - ▶ Use stored values to update weight and bias
- ▶ This is inefficient and could be improved
 - ▶ Requires storing full gradient matrices and revisiting them later; more storage used than necessary, affecting run times
 - ▶ Several loops means longer computational times
- ▶ We seek to interlace these steps to make the algorithm more efficient
 - ▶ For testing purposes, use data from the simple test case
 - ▶ Train with identical initial parameters and time how long each algorithm takes
 - ▶ Use networks of varying sizes

Interlacing Computation Analysis

- ▶ Use $(2,n,n,1)$ network for varying n up to $n = 40$, since we know this reliably converges for the simple test case data by previous analysis
- ▶ Report ratio of time for the original training method against the time for the interlaced training method (implemented via new '*train_interlaced*' function)



- ▶ All ratios reliably larger than 1
- ▶ Interlaced algorithm is significantly and consistently faster than initial implementation of the training algorithm
- ▶ Aligns with prediction